

In [1]:

```
import keras
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
from keras import backend as K
from keras.layers import Dense, Dropout, Flatten, BatchNormalization
from keras.layers import Conv2D, MaxPooling2D
import seaborn as sns
```

Using TensorFlow backend.

In [2]:

```
import matplotlib.pyplot as plt
import numpy as np
import time
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [3]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape", X_train[0].shape)
print("Number of training examples :", X_test.shape[0], "and each image is of shape", X_test[0].shape)
```

Number of training examples : 60000 and each image is of shape
(28, 28)

Number of training examples : 10000 and each image is of shape
(28, 28)

In [5]:

```
# if you observe the input shape its 2 dimensional vector  
# for each image we have a (28*28) vector  
# we will convert the (28*28) vector into single dimensional vector of 1 * 784  
  
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])  
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [6]:

```
# after converting the input images from 3d to 2d vectors  
  
print("Number of training examples :", X_train.shape[0], "and each image is of shape (784)")  
print("Number of test examples :", X_test.shape[0], "and each image is of shape (784)")
```

Number of training examples : 60000 and each image is of shape (784)

Number of test examples : 10000 and each image is of shape (784)

In [7]:

```
# An example data point  
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 17  
5 26 166 255  
 247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  3  
0 36 94 154  
 170 253 253 253 253 253 225 172 253 242 195 64  0  0  
0  0  0  0  
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 25  
3 251 93 82  
 82 56 39  0  0  0  0  0  0  0  0  0  0  0  0  
0 18 219 253  
 253 253 253 253 198 182 247 241  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0 80 156 107 253 253 205 1  
1  0 43 154  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 139 25  
3 190  2  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0 11 190 253 70  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  
0  0 35 241  
 225 160 108 1  0  0  0  0  0  0  0  0  0  0  0  
0  0  0  0  
  0  0  0  0  0  0  0  0  0  0 81 240 253 253 119 2  
5  0  0  0  
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

```

0  0  0  0
  0  0  45 186 253 253 150 27  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 16 9
3 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0 249 253 249 64  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 4
6 130 183 253
 253 207  2  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0 39 148 229 253 253 253 250 182  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 24 114 22
1 253 253 253
 253 201 78  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0 18 171 219 253 25
3 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
 55 172 226 253 253 253 253 244 133 11  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0 136 253 253 253 21
2 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]

```

In [8]:

```

# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize t
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

```

```

X_train = X_train/255
X_test = X_test/255

```

In [9]:

```
# example data point after normlizing
print(X_train[0])
```

```
[0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0.]
```

In [10]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0.
 1. 0. 0. 0. 0.]
```

In [11]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation

# some model parameters

batch_size = 128
num_classes = 10
epochs = 20

# input image dimensions
img_rows, img_cols = 28, 28
```

In [12]:

```
if K.image_data_format() == 'channels_first':
    x_train = X_train.reshape(X_train.shape[0], 1, img_rows, img_cols)
    x_test = X_test.reshape(X_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = X_train.reshape(X_train.shape[0], img_rows, img_cols, 1)
    x_test = X_test.reshape(X_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
```

In [13]:

```
x_train[0]
```

Out[13]:

[illegible]

**Model 1 : 2 Convolutional Layers with 3X3 Filters, 1
Dense layer (ReLU Activation) and Adam
Optimizer**

In [14]:

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                activation='relu',
                input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
print(model.summary())

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model.fit(x_train, Y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, Y_test))
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
flatten_1 (Flatten)	(None, 9216)	0
dense_1 (Dense)	(None, 128)	1179776
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290

Total params: 1,199,882
Trainable params: 1,199,882
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 82us/step -
loss: 0.2350 - accuracy: 0.9279 - val_loss: 0.0516 - val_accu-
acy: 0.9830

Epoch 2/20

60000/60000 [=====] - 3s 51us/step -
loss: 0.0848 - accuracy: 0.9741 - val_loss: 0.0365 - val_accu-
acy: 0.9879

Epoch 3/20

60000/60000 [=====] - 3s 52us/step -
loss: 0.0626 - accuracy: 0.9808 - val_loss: 0.0357 - val_accu-
acy: 0.9886

Epoch 4/20

60000/60000 [=====] - 3s 53us/step -
loss: 0.0514 - accuracy: 0.9841 - val_loss: 0.0336 - val_accu-
acy: 0.9878

Epoch 5/20

60000/60000 [=====] - 3s 52us/step -
loss: 0.0448 - accuracy: 0.9861 - val_loss: 0.0272 - val_accu-
acy: 0.9911

Epoch 6/20

60000/60000 [=====] - 3s 53us/step -
loss: 0.0387 - accuracy: 0.9876 - val_loss: 0.0258 - val_accu-
acy: 0.9909

Epoch 7/20

60000/60000 [=====] - 3s 52us/step -
loss: 0.0348 - accuracy: 0.9888 - val_loss: 0.0292 - val_accu-
acy: 0.9908

Epoch 8/20

60000/60000 [=====] - 3s 51us/step -
loss: 0.0313 - accuracy: 0.9898 - val_loss: 0.0262 - val_accu-
acy: 0.9920

Epoch 9/20

60000/60000 [=====] - 3s 51us/step -
loss: 0.0276 - accuracy: 0.9911 - val_loss: 0.0278 - val_accu-
acy: 0.9919

Epoch 10/20

60000/60000 [=====] - 3s 51us/step -
loss: 0.0260 - accuracy: 0.9917 - val_loss: 0.0284 - val_accu-
acy: 0.9920

Epoch 11/20

60000/60000 [=====] - 3s 53us/step -
loss: 0.0258 - accuracy: 0.9918 - val_loss: 0.0286 - val_accu-
acy: 0.9920

Epoch 12/20

60000/60000 [=====] - 3s 53us/step -

loss: 0.0223 - accuracy: 0.9927 - val_loss: 0.0278 - val_accuracy: 0.9924
Epoch 13/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0207 - accuracy: 0.9930 - val_loss: 0.0295 - val_accuracy: 0.9927
Epoch 14/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0177 - accuracy: 0.9940 - val_loss: 0.0298 - val_accuracy: 0.9926
Epoch 15/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0160 - accuracy: 0.9945 - val_loss: 0.0301 - val_accuracy: 0.9918
Epoch 16/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0190 - accuracy: 0.9938 - val_loss: 0.0306 - val_accuracy: 0.9929
Epoch 17/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0170 - accuracy: 0.9942 - val_loss: 0.0274 - val_accuracy: 0.9932
Epoch 18/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0171 - accuracy: 0.9942 - val_loss: 0.0274 - val_accuracy: 0.9931
Epoch 19/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0142 - accuracy: 0.9955 - val_loss: 0.0280 - val_accuracy: 0.9927
Epoch 20/20
60000/60000 [=====] - 3s 52us/step -
loss: 0.0118 - accuracy: 0.9962 - val_loss: 0.0275 - val_accuracy: 0.9936

In [15]:

```
score = model.evaluate(x_test, Y_test, verbose=0)
score1 = score[0]
accuracy1 = score[1]
print('Test score:', score1)
print('Test accuracy:', accuracy1)

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, epochs+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb


# we will get val_loss and val_acc only when you pass the paramter validation
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to nb

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.027506851389828297

Test accuracy: 0.9936000108718872



Model 2 : 2 Convolutional Layers with 5X5 Filters, 2 Dense layers (ReLU Activation) and Adam Optimizer

In [16]:

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(5,5), padding='same', activation='relu', in
model.add(Conv2D(64, (5,5), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
print(model.summary())

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model.fit(x_train, Y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, Y_test))
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
===		
conv2d_3 (Conv2D)	(None, 28, 28, 32)	832
conv2d_4 (Conv2D)	(None, 24, 24, 64)	51264
max_pooling2d_2 (MaxPooling2	(None, 12, 12, 64)	0
dropout_3 (Dropout)	(None, 12, 12, 64)	0
flatten_2 (Flatten)	(None, 9216)	0
dense_3 (Dense)	(None, 256)	2359552
dropout_4 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 256)	65792

batch_normalization_1 (Batch Normalization)	(None, 256)	1024
---	-------------	------

dropout_5 (Dropout)	(None, 256)	0
---------------------	-------------	---

dense_5 (Dense)	(None, 10)	2570
-----------------	------------	------

=====
===

Total params: 2,481,034

Trainable params: 2,480,522

Non-trainable params: 512

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 77us/step -
loss: 0.2347 - accuracy: 0.9287 - val_loss: 0.0455 - val_accu-
acy: 0.9859

Epoch 2/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0744 - accuracy: 0.9793 - val_loss: 0.0397 - val_accu-
acy: 0.9880

Epoch 3/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0566 - accuracy: 0.9827 - val_loss: 0.0290 - val_accu-
acy: 0.9909

Epoch 4/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0463 - accuracy: 0.9863 - val_loss: 0.0272 - val_accu-
acy: 0.9916

Epoch 5/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0362 - accuracy: 0.9892 - val_loss: 0.0209 - val_accu-
acy: 0.9935

Epoch 6/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0318 - accuracy: 0.9905 - val_loss: 0.0259 - val_accu-
acy: 0.9920

Epoch 7/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0282 - accuracy: 0.9913 - val_loss: 0.0290 - val_accu-
acy: 0.9920

Epoch 8/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0238 - accuracy: 0.9925 - val_loss: 0.0280 - val_accu-
acy: 0.9922

Epoch 9/20

60000/60000 [=====] - 4s 67us/step -
loss: 0.0230 - accuracy: 0.9933 - val_loss: 0.0249 - val_accu-
acy: 0.9924

Epoch 10/20
60000/60000 [=====] - 4s 66us/step -
loss: 0.0221 - accuracy: 0.9932 - val_loss: 0.0246 - val_accu-
racy: 0.9921
Epoch 11/20
60000/60000 [=====] - 4s 67us/step -
loss: 0.0200 - accuracy: 0.9941 - val_loss: 0.0291 - val_accu-
racy: 0.9934
Epoch 12/20
60000/60000 [=====] - 4s 67us/step -
loss: 0.0202 - accuracy: 0.9937 - val_loss: 0.0243 - val_accu-
racy: 0.9939
Epoch 13/20
60000/60000 [=====] - 4s 67us/step -
loss: 0.0170 - accuracy: 0.9947 - val_loss: 0.0259 - val_accu-
racy: 0.9931
Epoch 14/20
60000/60000 [=====] - 4s 67us/step -
loss: 0.0166 - accuracy: 0.9947 - val_loss: 0.0243 - val_accu-
racy: 0.9932
Epoch 15/20
60000/60000 [=====] - 4s 66us/step -
loss: 0.0133 - accuracy: 0.9963 - val_loss: 0.0258 - val_accu-
racy: 0.9938
Epoch 16/20
60000/60000 [=====] - 4s 67us/step -
loss: 0.0146 - accuracy: 0.9955 - val_loss: 0.0257 - val_accu-
racy: 0.9929
Epoch 17/20
60000/60000 [=====] - 4s 66us/step -
loss: 0.0137 - accuracy: 0.9960 - val_loss: 0.0251 - val_accu-
racy: 0.9941
Epoch 18/20
60000/60000 [=====] - 4s 66us/step -
loss: 0.0121 - accuracy: 0.9962 - val_loss: 0.0231 - val_accu-
racy: 0.9942
Epoch 19/20
60000/60000 [=====] - 4s 66us/step -
loss: 0.0129 - accuracy: 0.9958 - val_loss: 0.0272 - val_accu-
racy: 0.9934
Epoch 20/20
60000/60000 [=====] - 4s 67us/step -
loss: 0.0100 - accuracy: 0.9970 - val_loss: 0.0269 - val_accu-
racy: 0.9936

In [17]:

```
score = model.evaluate(x_test, Y_test, verbose=0)
score2 = score[0]
accuracy2 = score[1]
print('Test score:', score2)
print('Test accuracy:', accuracy2)


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0268641913799965

Test accuracy: 0.9936000108718872



Model 3 : 2 Convolutional Layers with 11X11 Filters, 2 Dense layer (ReLU Activation) and Adam Optimizer

In [18]:

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(11, 11), padding='same', activation='relu'))
model.add(Conv2D(64, (11, 11), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
print(model.summary())

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['acc'])
history = model.fit(x_train, Y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, Y_test))
```

Model: "sequential_3"

Layer (type) m #	Output Shape	Para
=====		
conv2d_5 (Conv2D)	(None, 28, 28, 32)	3904
=====		
conv2d_6 (Conv2D) 72	(None, 18, 18, 64)	2478
=====		
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 64)	0
=====		
dropout_6 (Dropout)	(None, 9, 9, 64)	0
=====		
flatten_3 (Flatten)	(None, 5184)	0
=====		
dense_6 (Dense) 360	(None, 256)	1327
=====		
dropout_7 (Dropout)	(None, 256)	0

dense_7 (Dense)	(None, 256)	6579
2		
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
dropout_8 (Dropout)	(None, 256)	0
dense_8 (Dense)	(None, 10)	2570
=====		
=====		
Total params: 1,648,522		
Trainable params: 1,648,010		
Non-trainable params: 512		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 83us/step
- loss: 0.2395 - accuracy: 0.9278 - val_loss: 0.0454 - val_accuracy: 0.9848

Epoch 2/20

60000/60000 [=====] - 4s 72us/step
- loss: 0.0704 - accuracy: 0.9804 - val_loss: 0.0269 - val_accuracy: 0.9911

Epoch 3/20

60000/60000 [=====] - 4s 72us/step
- loss: 0.0496 - accuracy: 0.9856 - val_loss: 0.0245 - val_accuracy: 0.9924

Epoch 4/20

60000/60000 [=====] - 4s 72us/step
- loss: 0.0392 - accuracy: 0.9884 - val_loss: 0.0266 - val_accuracy: 0.9922

Epoch 5/20

60000/60000 [=====] - 4s 72us/step
- loss: 0.0368 - accuracy: 0.9894 - val_loss: 0.0240 - val_accuracy: 0.9929

Epoch 6/20

60000/60000 [=====] - 4s 72us/step
- loss: 0.0285 - accuracy: 0.9918 - val_loss: 0.0226 - val_accuracy: 0.9935

Epoch 7/20

60000/60000 [=====] - 4s 72us/step
- loss: 0.0250 - accuracy: 0.9925 - val_loss: 0.0247 - val_accuracy: 0.9923

Epoch 8/20

60000/60000 [=====] - 4s 71us/step
- loss: 0.0243 - accuracy: 0.9928 - val_loss: 0.0349 - val_

accuracy: 0.9904
Epoch 9/20
60000/60000 [=====] - 4s 71us/step
- loss: 0.0232 - accuracy: 0.9931 - val_loss: 0.0308 - val_
accuracy: 0.9918
Epoch 10/20
60000/60000 [=====] - 4s 72us/step
- loss: 0.0203 - accuracy: 0.9940 - val_loss: 0.0208 - val_
accuracy: 0.9949
Epoch 11/20
60000/60000 [=====] - 4s 72us/step
- loss: 0.0171 - accuracy: 0.9950 - val_loss: 0.0277 - val_
accuracy: 0.9925
Epoch 12/20
60000/60000 [=====] - 4s 74us/step
- loss: 0.0159 - accuracy: 0.9951 - val_loss: 0.0235 - val_
accuracy: 0.9935
Epoch 13/20
60000/60000 [=====] - 4s 73us/step
- loss: 0.0160 - accuracy: 0.9952 - val_loss: 0.0295 - val_
accuracy: 0.9930
Epoch 14/20
60000/60000 [=====] - 4s 71us/step
- loss: 0.0152 - accuracy: 0.9952 - val_loss: 0.0243 - val_
accuracy: 0.9933
Epoch 15/20
60000/60000 [=====] - 4s 72us/step
- loss: 0.0114 - accuracy: 0.9964 - val_loss: 0.0228 - val_
accuracy: 0.9930
Epoch 16/20
60000/60000 [=====] - 4s 71us/step
- loss: 0.0119 - accuracy: 0.9965 - val_loss: 0.0217 - val_
accuracy: 0.9936
Epoch 17/20
60000/60000 [=====] - 4s 71us/step
- loss: 0.0098 - accuracy: 0.9971 - val_loss: 0.0251 - val_
accuracy: 0.9935
Epoch 18/20
60000/60000 [=====] - 4s 71us/step
- loss: 0.0089 - accuracy: 0.9973 - val_loss: 0.0223 - val_
accuracy: 0.9940
Epoch 19/20
60000/60000 [=====] - 4s 71us/step
- loss: 0.0091 - accuracy: 0.9969 - val_loss: 0.0283 - val_
accuracy: 0.9934
Epoch 20/20
60000/60000 [=====] - 4s 71us/step
- loss: 0.0111 - accuracy: 0.9966 - val_loss: 0.0250 - val_
accuracy: 0.9942

In [19]:


```
score = model.evaluate(x_test, Y_test, verbose=0)
score3 = score[0]
accuracy3 = score[1]
print('Test score:', score3)
print('Test accuracy:', accuracy3)

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.025038751071174595
Test accuracy: 0.9941999912261963



Model 4 : 2 Convolutional Layers with 11X11 Filters (No Padding), 2 Dense layer (ReLU Activation) and Adam Optimizer

In [20]:

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(11, 11), activation='relu', input_shape=inp
model.add(Conv2D(64, (11, 11), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
print(model.summary())

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model.fit(x_train, Y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, Y_test))
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 18, 18, 32)	3904
conv2d_8 (Conv2D)	(None, 8, 8, 64)	2472
max_pooling2d_4 (MaxPooling2D)	(None, 4, 4, 64)	0
dropout_9 (Dropout)	(None, 4, 4, 64)	0
flatten_4 (Flatten)	(None, 1024)	0
dense_9 (Dense)	(None, 256)	262400
dropout_10 (Dropout)	(None, 256)	0

dense_10 (Dense)	(None, 256)	6579
2		
batch_normalization_3 (Batch Normalization)	(None, 256)	1024
dropout_11 (Dropout)	(None, 256)	0
dense_11 (Dense)	(None, 10)	2570
=====		
=====		
Total params: 583,562		
Trainable params: 583,050		
Non-trainable params: 512		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 62us/step
- loss: 0.2960 - accuracy: 0.9097 - val_loss: 0.0492 - val_accuracy: 0.9846

Epoch 2/20

60000/60000 [=====] - 3s 52us/step
- loss: 0.0806 - accuracy: 0.9767 - val_loss: 0.0536 - val_accuracy: 0.9844

Epoch 3/20

60000/60000 [=====] - 3s 53us/step
- loss: 0.0604 - accuracy: 0.9830 - val_loss: 0.0394 - val_accuracy: 0.9896

Epoch 4/20

60000/60000 [=====] - 3s 52us/step
- loss: 0.0482 - accuracy: 0.9865 - val_loss: 0.0317 - val_accuracy: 0.9905

Epoch 5/20

60000/60000 [=====] - 3s 53us/step
- loss: 0.0412 - accuracy: 0.9880 - val_loss: 0.0284 - val_accuracy: 0.9927

Epoch 6/20

60000/60000 [=====] - 3s 54us/step
- loss: 0.0352 - accuracy: 0.9898 - val_loss: 0.0324 - val_accuracy: 0.9907

Epoch 7/20

60000/60000 [=====] - 3s 52us/step
- loss: 0.0332 - accuracy: 0.9901 - val_loss: 0.0360 - val_accuracy: 0.9901

Epoch 8/20

60000/60000 [=====] - 3s 54us/step
- loss: 0.0290 - accuracy: 0.9918 - val_loss: 0.0314 - val_accuracy: 0.9908

Epoch 9/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0281 - accuracy: 0.9915 - val_loss: 0.0295 - val_
accuracy: 0.9926
Epoch 10/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0229 - accuracy: 0.9936 - val_loss: 0.0312 - val_
accuracy: 0.9922
Epoch 11/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0212 - accuracy: 0.9936 - val_loss: 0.0392 - val_
accuracy: 0.9909
Epoch 12/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0231 - accuracy: 0.9930 - val_loss: 0.0344 - val_
accuracy: 0.9915
Epoch 13/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0183 - accuracy: 0.9945 - val_loss: 0.0375 - val_
accuracy: 0.9906
Epoch 14/20
60000/60000 [=====] - 3s 53us/step
- loss: 0.0164 - accuracy: 0.9952 - val_loss: 0.0295 - val_
accuracy: 0.9920
Epoch 15/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0158 - accuracy: 0.9954 - val_loss: 0.0374 - val_
accuracy: 0.9904
Epoch 16/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0166 - accuracy: 0.9951 - val_loss: 0.0392 - val_
accuracy: 0.9897
Epoch 17/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0159 - accuracy: 0.9952 - val_loss: 0.0335 - val_
accuracy: 0.9928
Epoch 18/20
60000/60000 [=====] - 3s 55us/step
- loss: 0.0125 - accuracy: 0.9962 - val_loss: 0.0375 - val_
accuracy: 0.9921
Epoch 19/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0141 - accuracy: 0.9959 - val_loss: 0.0352 - val_
accuracy: 0.9916
Epoch 20/20
60000/60000 [=====] - 3s 54us/step
- loss: 0.0107 - accuracy: 0.9967 - val_loss: 0.0354 - val_
accuracy: 0.9924

In [21]:

```
score = model.evaluate(x_test, Y_test, verbose=0)
score4 = score[0]
accuracy4 = score[1]
print('Test score:', score4)
print('Test accuracy:', accuracy4)


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0353821564207712

Test accuracy: 0.9923999905586243



Model 5 : 3 Convolutional Layers with 11X11 Filters, 2 Dense layer (ReLU Activation), 2 MaxPool and Adam Optimizer

In [22]:

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(11,11), padding='same', activation='relu',
model.add(Conv2D(64, (11, 11), padding='same', activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (11, 11), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
print(model.summary())

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['ac
history = model.fit(x_train, Y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, Y_test))
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 28, 28, 32)	3904
conv2d_10 (Conv2D)	(None, 28, 28, 64)	24784
max_pooling2d_5 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_11 (Conv2D)	(None, 4, 4, 64)	49568
max_pooling2d_6 (MaxPooling2D)	(None, 2, 2, 64)	0
dropout_12 (Dropout)	(None, 2, 2, 64)	0

flatten_5 (Flatten)	(None, 256)	0
dense_12 (Dense) 2	(None, 256)	6579
dropout_13 (Dropout)	(None, 256)	0
dense_13 (Dense) 2	(None, 256)	6579
batch_normalization_4 (Batch Normalization)	(None, 256)	1024
dropout_14 (Dropout)	(None, 256)	0
dense_14 (Dense)	(None, 10)	2570
=====		
=====		
Total params: 882,634		
Trainable params: 882,122		
Non-trainable params: 512		

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 120us/step - loss: 0.2782 - accuracy: 0.9136 - val_loss: 0.0603 - val_accuracy: 0.9805

Epoch 2/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0703 - accuracy: 0.9808 - val_loss: 0.0282 - val_accuracy: 0.9919

Epoch 3/20

60000/60000 [=====] - 7s 108us/step - loss: 0.0471 - accuracy: 0.9869 - val_loss: 0.0236 - val_accuracy: 0.9923

Epoch 4/20

60000/60000 [=====] - 7s 109us/step - loss: 0.0407 - accuracy: 0.9891 - val_loss: 0.0485 - val_accuracy: 0.9878

Epoch 5/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0351 - accuracy: 0.9904 - val_loss: 0.0258 - val_accuracy: 0.9927

Epoch 6/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0279 - accuracy: 0.9921 - val_loss: 0.0248 - va

l_accuracy: 0.9930
Epoch 7/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0240 - accuracy: 0.9932 - val_loss: 0.0245 - val_accuracy: 0.9928
Epoch 8/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0213 - accuracy: 0.9942 - val_loss: 0.0294 - val_accuracy: 0.9925
Epoch 9/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0207 - accuracy: 0.9940 - val_loss: 0.0283 - val_accuracy: 0.9928
Epoch 10/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0183 - accuracy: 0.9948 - val_loss: 0.0263 - val_accuracy: 0.9926
Epoch 11/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0144 - accuracy: 0.9960 - val_loss: 0.0355 - val_accuracy: 0.9903
Epoch 12/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0170 - accuracy: 0.9953 - val_loss: 0.0307 - val_accuracy: 0.9925
Epoch 13/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0151 - accuracy: 0.9958 - val_loss: 0.0263 - val_accuracy: 0.9928
Epoch 14/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0118 - accuracy: 0.9967 - val_loss: 0.0261 - val_accuracy: 0.9937
Epoch 15/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0123 - accuracy: 0.9966 - val_loss: 0.0217 - val_accuracy: 0.9936
Epoch 16/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0088 - accuracy: 0.9976 - val_loss: 0.0214 - val_accuracy: 0.9945
Epoch 17/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0120 - accuracy: 0.9968 - val_loss: 0.0346 - val_accuracy: 0.9924
Epoch 18/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0108 - accuracy: 0.9967 - val_loss: 0.0234 - val_accuracy: 0.9940
Epoch 19/20
60000/60000 [=====] - 6s 106us/step - loss: 0.0092 - accuracy: 0.9976 - val_loss: 0.0325 - val_accuracy: 0.9925

Epoch 20/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0078 - accuracy: 0.9979 - val_loss: 0.0280 - val_accuracy: 0.9934

In [23]:

```
score = model.evaluate(x_test, Y_test, verbose=0)
score5 = score[0]
accuracy5 = score[1]
print('Test score:', score5)
print('Test accuracy:', accuracy5)


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, epochs+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.028024690563664523

Test accuracy: 0.993399977684021




Vizualizing Results

In [24]:

```
import seaborn as sns
models = ['2 Conv, 1 Dense 3X3 Filter', '2 Conv, 2 Dense 5X5 Filter', '2 Conv, 2 Dense 11X11 Filter', '2 Conv, 2 Dense 11X11 Filter No Padding', '3 Conv, 2 Dense 11X11 Filter']
scores = [score1, score2, score3, score4, score5]
accuracies = [accuracy1, accuracy2, accuracy3, accuracy4, accuracy5]
```

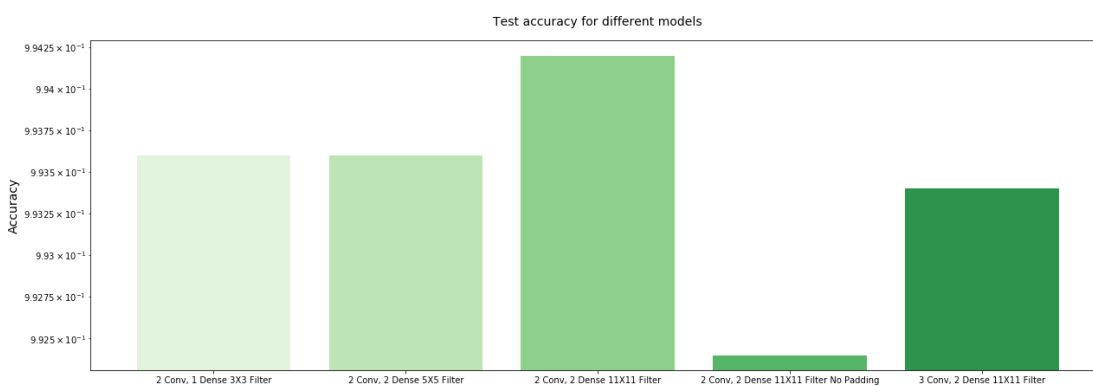
In [25]:

```
plt.figure(figsize=(21,7))
fig = plt.bar(models, scores, color=sns.color_palette("Reds",6))
plt.ylabel('Loss', size = 14)
plt.yscale('log')
plt.title('Test loss for different models', size = 14, y = 1.03)
plt.show()
```



In [26]:

```
plt.figure(figsize=(21,7))
fig = plt.bar(models, accuracies, color=sns.color_palette("Greens",6))
plt.ylabel('Accuracy', size = 14)
plt.yscale('log')
plt.title('Test accuracy for different models', size = 14, y = 1.03)
plt.show()
```



Conclusion

In [27]:

```
from prettytable import PrettyTable

table = PrettyTable()
table.field_names = ["Model", "Loss", "Accuracy"]

table.add_row([models[0], score1, accuracy1])
table.add_row([models[1], score2, accuracy2])
table.add_row([models[2], score3, accuracy3])
table.add_row([models[3], score4, accuracy4])
table.add_row([models[4], score5, accuracy5])
print(table)
```

```
+-----+-----+
---+-----+
|          Model          |          Loss          |
|          Accuracy       |                         |
+-----+-----+
---+-----+
|      2 Conv, 1 Dense 3X3 Filter      | 0.0275068513898282 |
97 | 0.9936000108718872 |
|      2 Conv, 2 Dense 5X5 Filter      | 0.026864191379996 |
5  | 0.9936000108718872 |
|      2 Conv, 2 Dense 11X11 Filter     | 0.0250387510711745 |
95 | 0.9941999912261963 |
| 2 Conv, 2 Dense 11X11 Filter No Padding | 0.035382156420771 |
2  | 0.9923999905586243 |
|      3 Conv, 2 Dense 11X11 Filter     | 0.0280246905636645 |
23 | 0.993399977684021  |
+-----+-----+
---+-----+
```